



Grid A — metrics x query types, model lines vs corpus size

Jina_v3 Qwen0.6b Qwen8b BGE_M3 Nomic_Embed_v2 Snowflake_v2 Promptriever GritLM

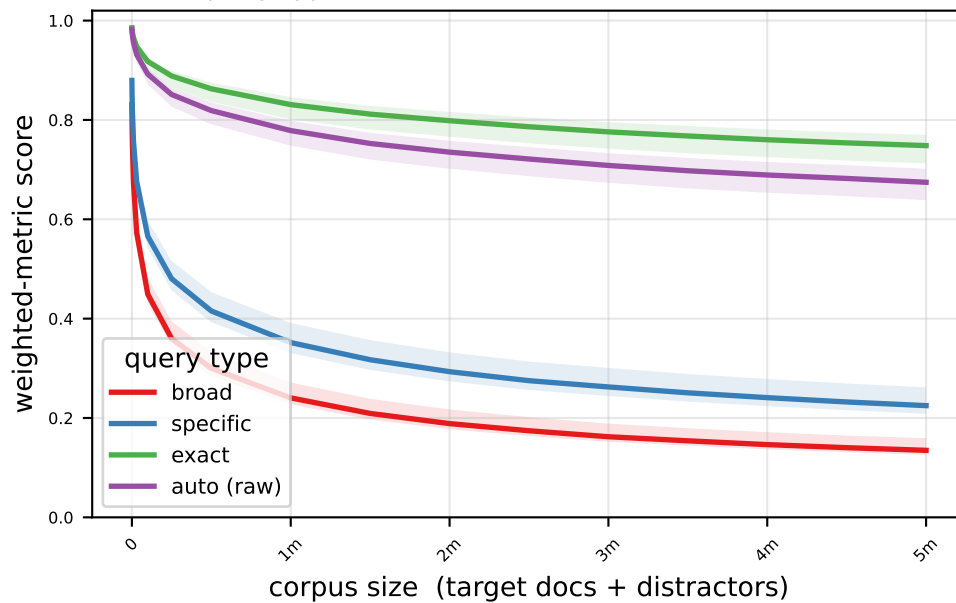


Grid B — models x query types, metric lines vs corpus size

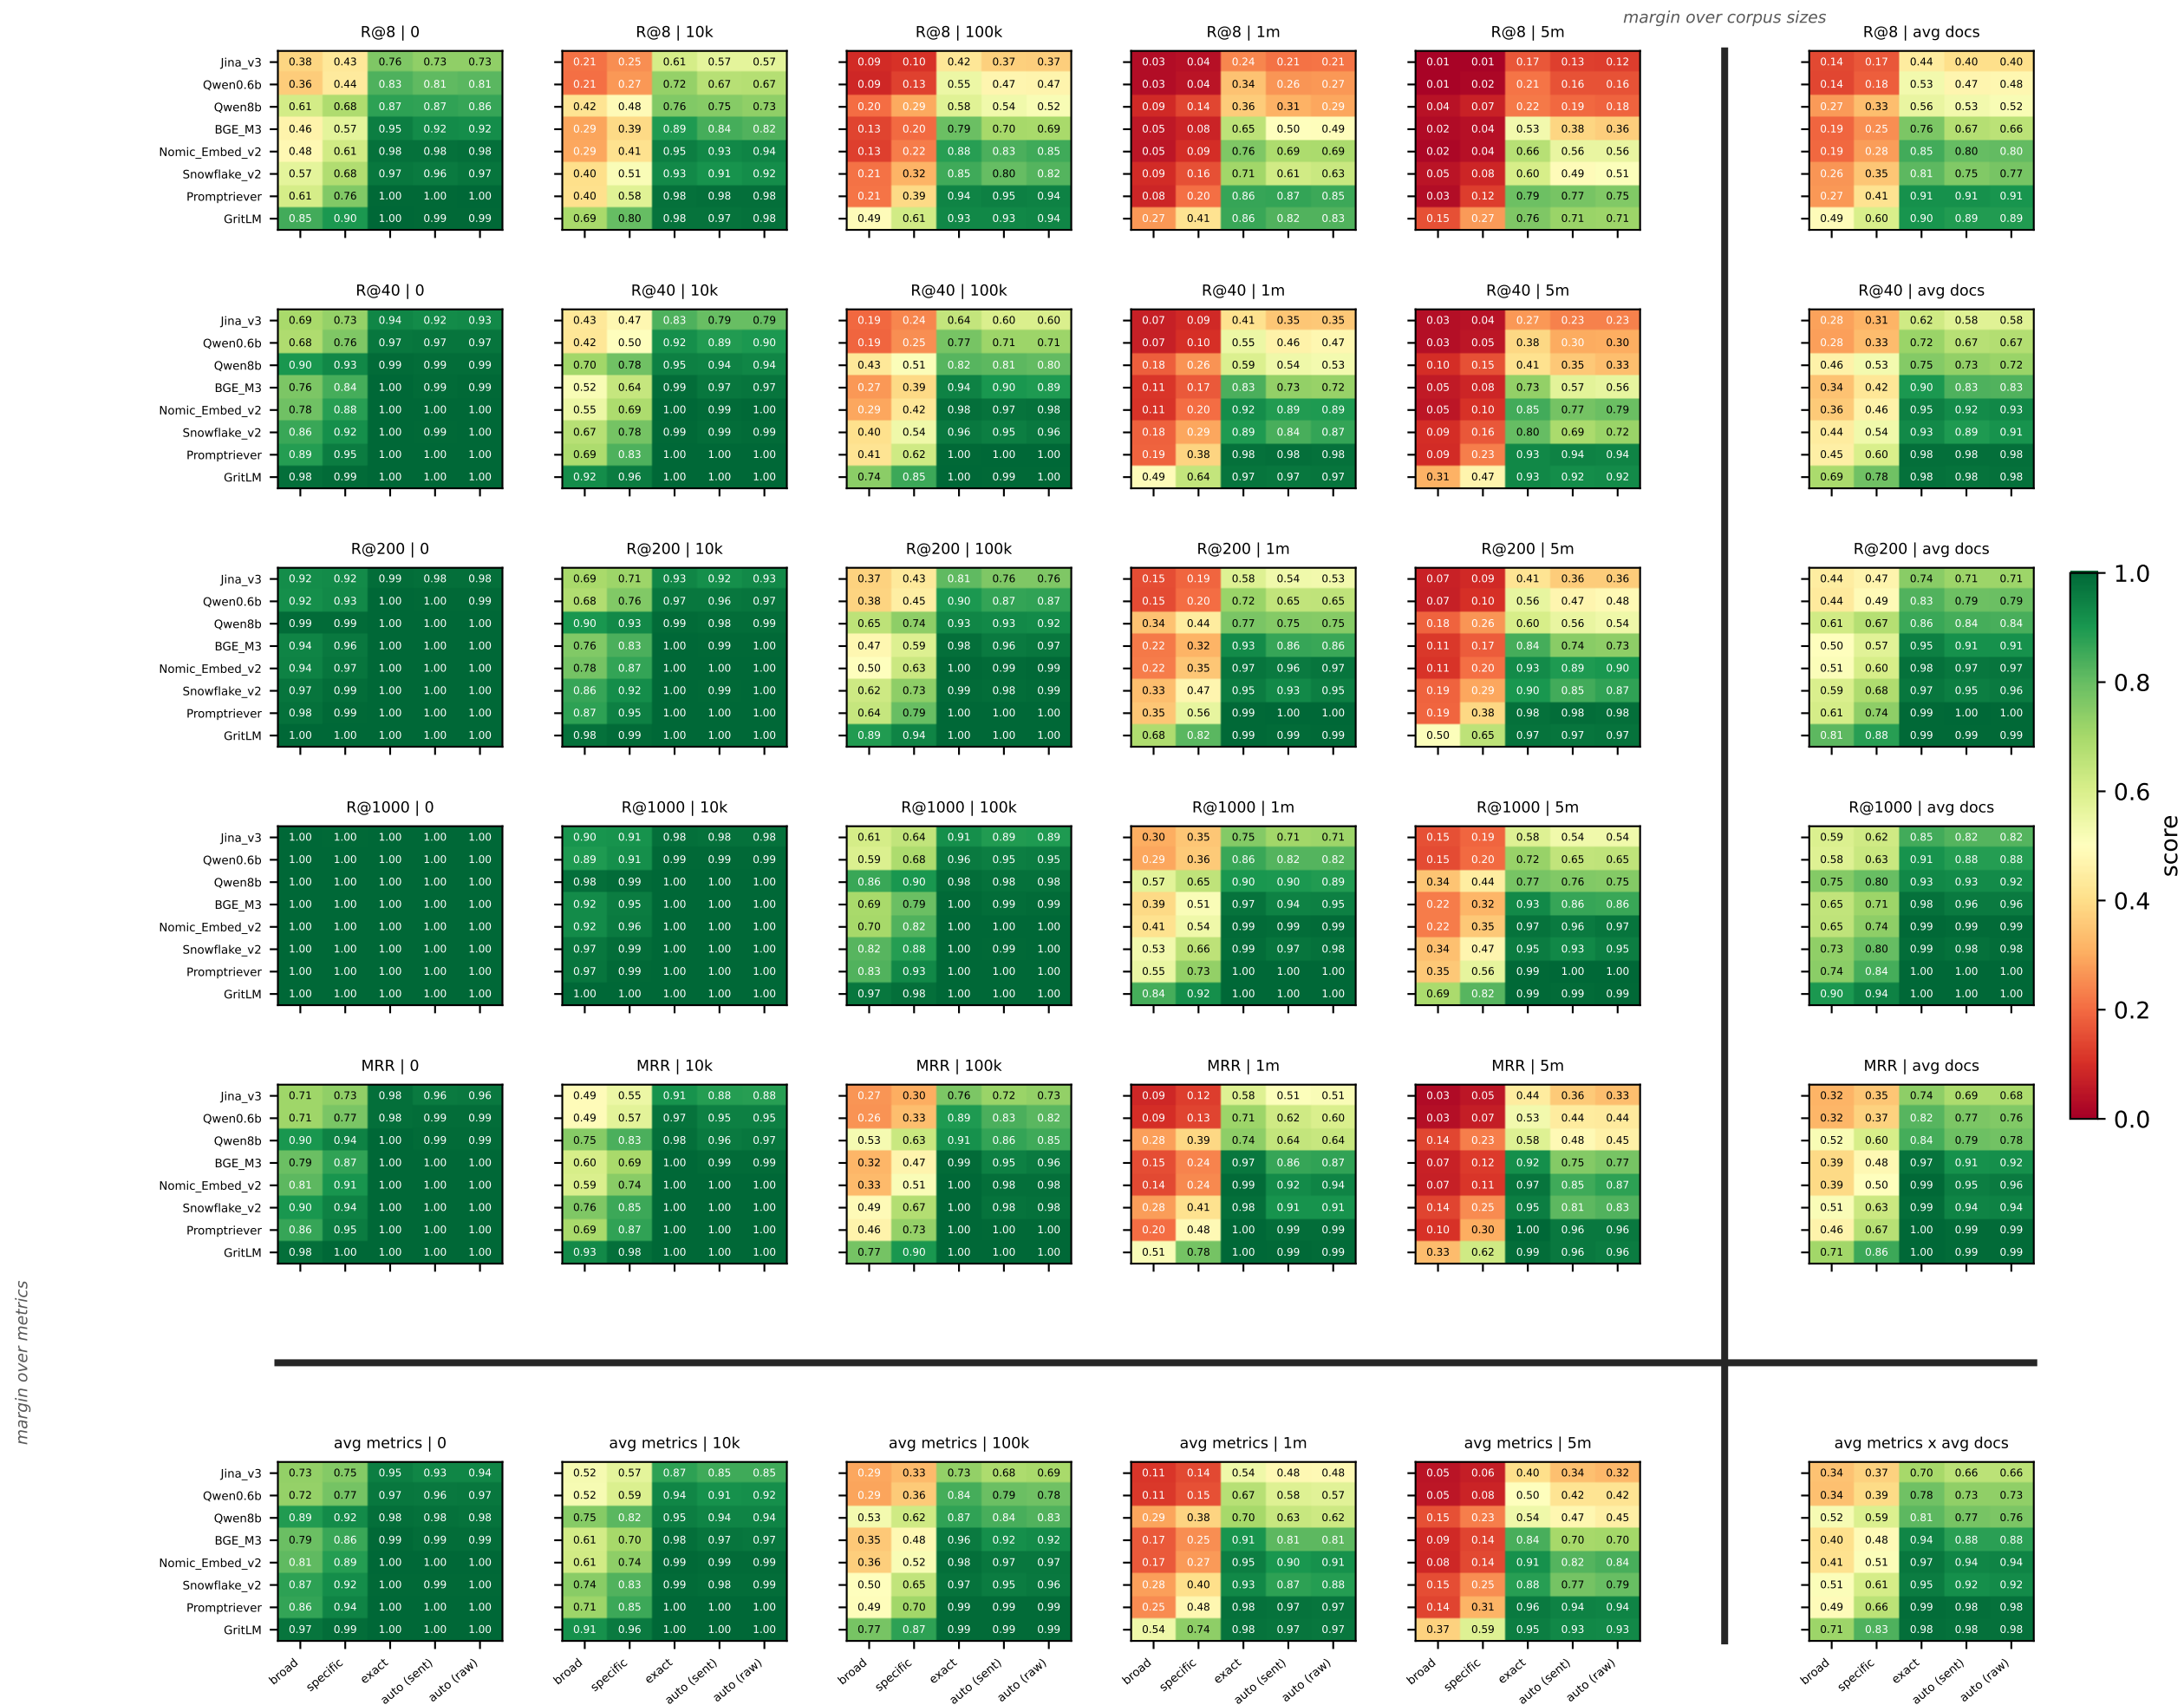
R@8 R@40 R@200 R@1000 MRR



Plot C — query-type bands (width = 0.10× min/max across models)

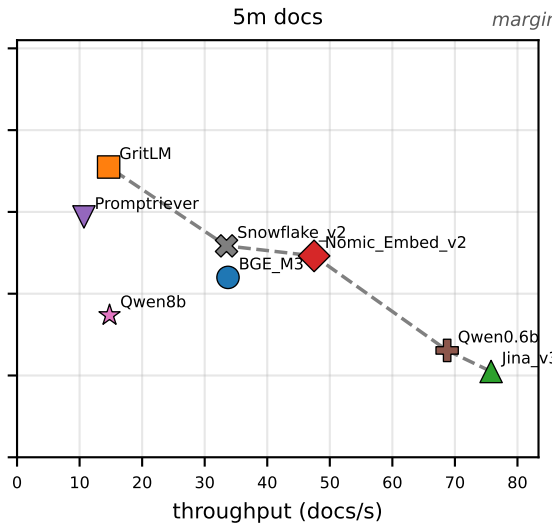
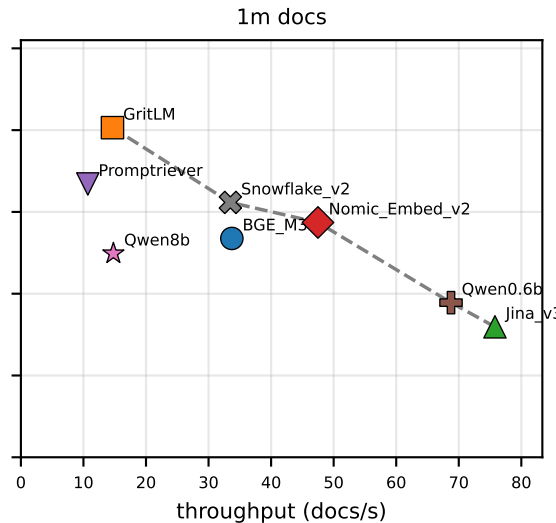
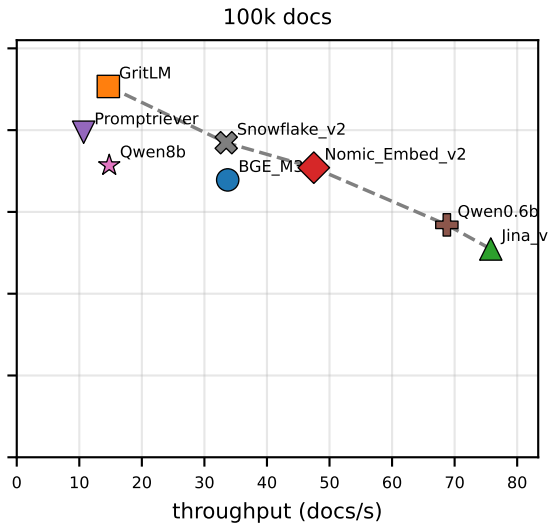
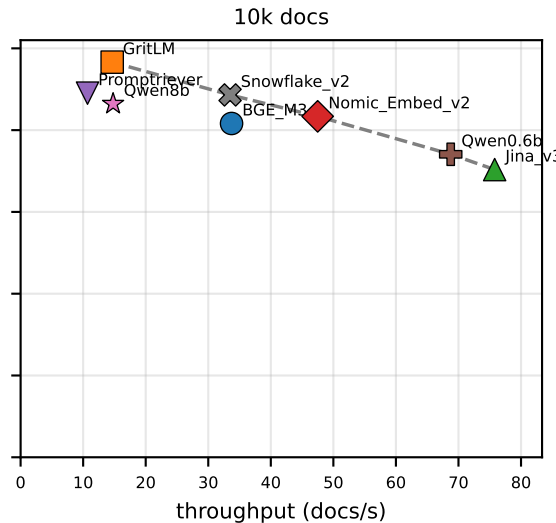
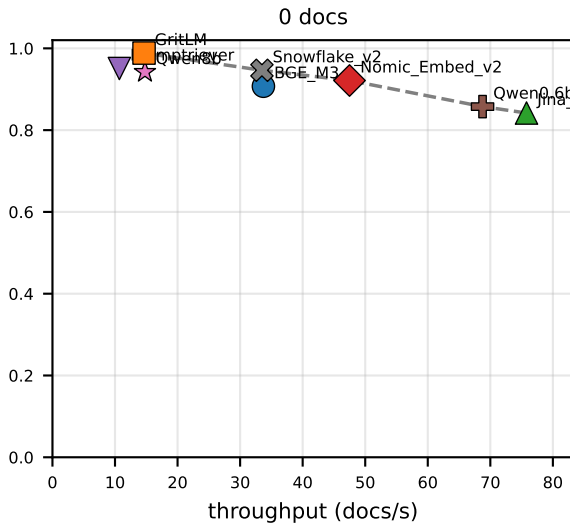


Grid D — metrics x corpus size; each cell = model x query-type heatmap



Pareto — throughput vs score per corpus size (dashed = frontier)

weighted score (queries x metrics)



margin over docs

